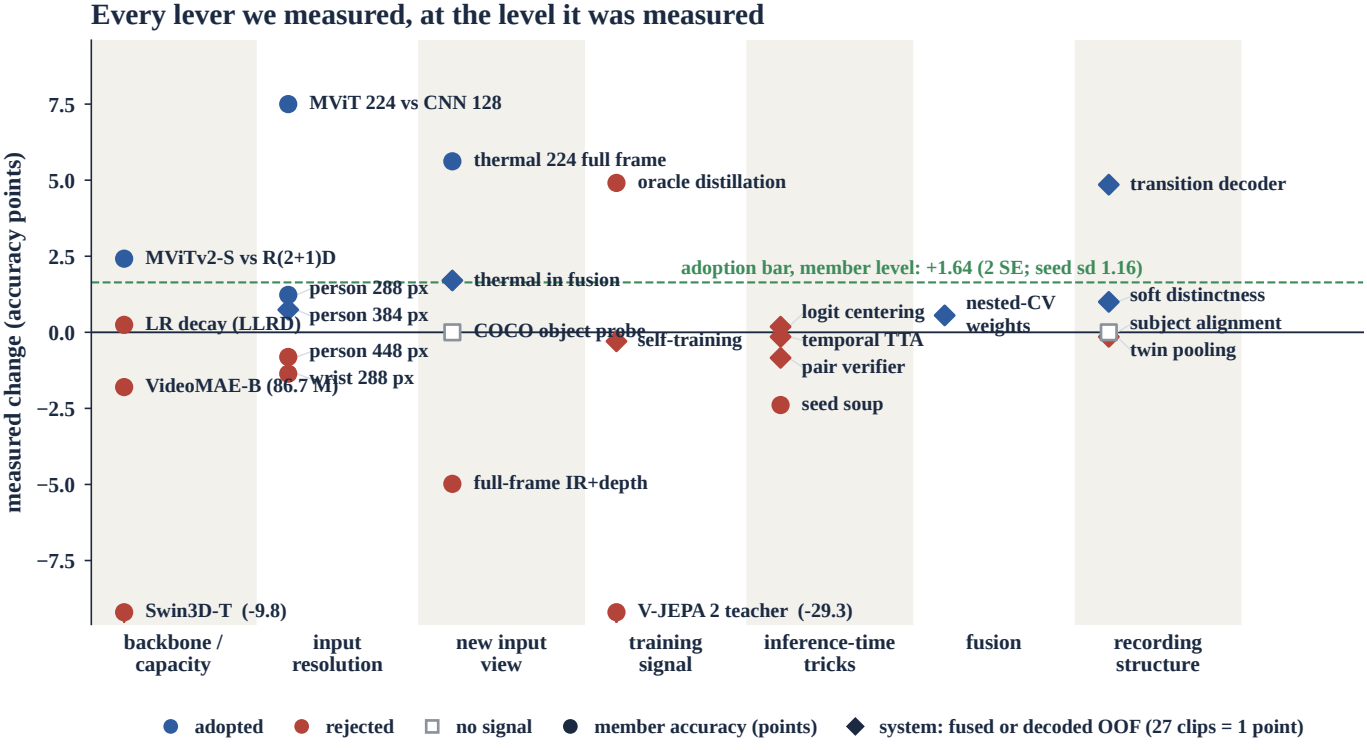
Every lever we measured, at the level it was measured
measured change (accuracy points)
7.5
5.0
2.5
0.0
−2.5
−5.0
−7.5
MViT 224 vs CNN 128
thermal 224 full frame
oracle distillation
transition decoder
MViTv2-S vs R(2+1)D
adoption bar, member level: +1.64 (2 SE; seed sd 1.16)
thermal in fusion
person 288 px
person 384 px
LR decay (LLRD)
COCO object probe
self-training
logit centering
temporal TTA
pair verifier
nested-CV weights
soft distinctness
subject alignment
twin pooling
person 448 px
wrist 288 px
VideoMAE-B (86.7 M)
seed soup
full-frame IR+depth
Swin3D-T (-9.8)
V-JEPA 2 teacher (-29.3)
backbone / capacity
input resolution
new input view
training signal
inference-time tricks
fusion
recording structure
adopted
rejected
no signal
member accuracy (points)
system: fused or decoded OOF (27 clips = 1 point)