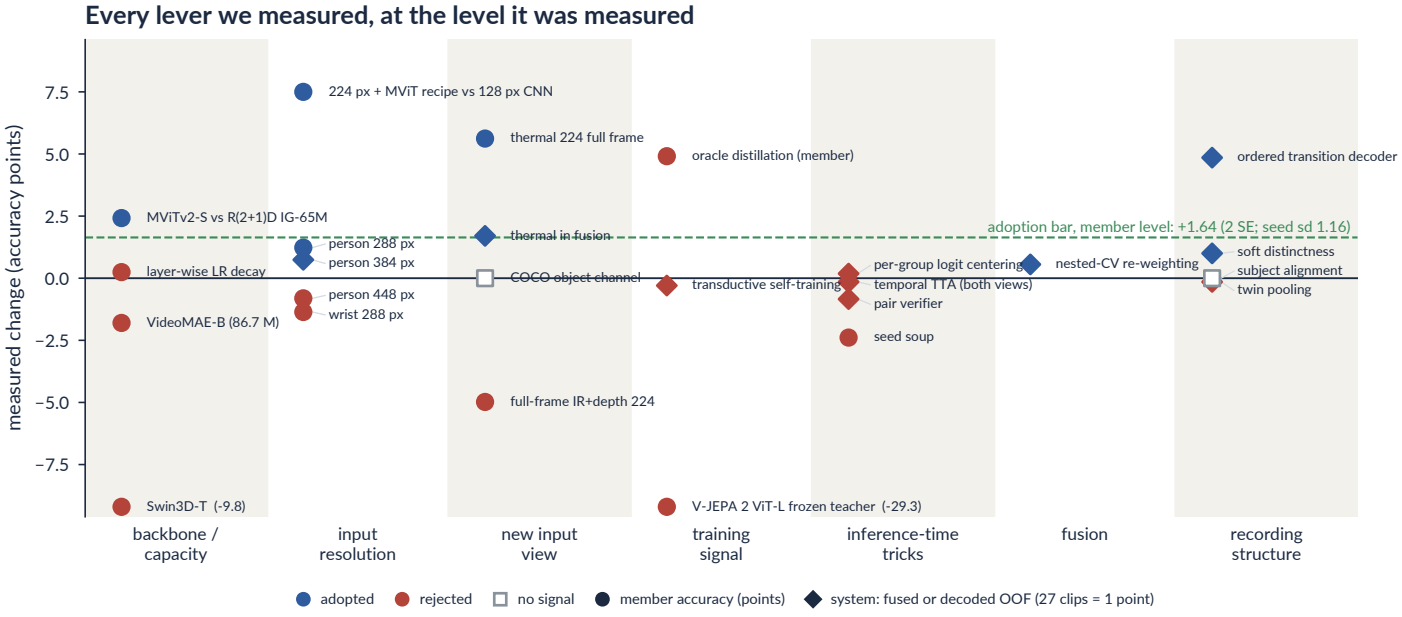
Every lever we measured, at the level it was measured
measured change (accuracy points)
7.5
5.0
2.5
0.0
−2.5
−5.0
−7.5
224 px + MViT recipe vs 128 px CNN
thermal 224 full frame
oracle distillation (member)
ordered transition decoder
MViTv2-S vs R(2+1)D IG-65M
adoption bar, member level: +1.64 (2 SE; seed sd 1.16)
person 288 px
person 384 px
thermal in fusion
soft distinctness
layer-wise LR decay
COCO object channel
transductive self-training
per-group logit centering
nested-CV re-weighting
subject alignment
temporal TTA (both views)
twin pooling
person 448 px
wrist 288 px
pair verifier
VideoMAE-B (86.7 M)
seed soup
full-frame IR+depth 224
Swin3D-T (-9.8)
V-JEPA 2 ViT-L frozen teacher (-29.3)
backbone / capacity
input resolution
new input view
training signal
inference-time tricks
fusion
recording structure
adopted
rejected
no signal
member accuracy (points)
system: fused or decoded OOF (27 clips = 1 point)