WHAT WE TRIED

every lever measured against the same bar: paired folds, 2 SE (seed sd 1.16 points)

MORE CAPACITY

VideoMAE-B -1.8 (0/4 folds) · Swin3D-T -9.8 · V-JEPA 2 teacher -29

MORE TRAINING TRICKS

distillation: student 157 vs 167 public · self-training: member +1.4, fusion -2 · LR decay +0.2 (noise)

MORE TEST-TIME TRICKS

temporal TTA -4 fused · seed soup -2.4 (4/4 folds) · pair verifier -0.8 · logit centering +5 clips (unreadable)

MORE INPUT RESOLUTION

person crop 224 to 384 px +1.9 member, +20 fused clips (448 px: -0.8, rejected)

A NEW SENSOR VIEW

thermal at 224 px full frame +5.6 member, +46 decoded clips, +5 public

BETTER FUSION

prior-corrected log-linear space (wrong space costs 4 points) · nested-CV weights +15 clips

RECORDING STRUCTURE

ordered transition decoder +131 decoded clips (+55 distinctness, +76 order) · +3 to +5 public

the bottleneck is not capacity: the residual error is subject-correlated (next slide: every lever placed on one axis)